

Ryan Yee - CS 760 Final Note Sheet

Information Theory

Entropy: $H(Y) = - \sum_{y \in Y} \mathbb{P}(y) \log_2(\mathbb{P}(y))$

$$H(Y|X) = \sum_{x \in X} \mathbb{P}(X=x) H(Y|X=x)$$

Information Gain: $H(Y) - H(Y|X)$

Gain Ratio: $\frac{H(Y) - H(Y|S)}{H(S)}$

Evaluation

Accuracy: $\frac{TP+TN}{TP+FP+TN+FN}$, Error: $1 - \text{Accuracy}$

Precision: $\frac{TP}{TP+FP}$, Recall: $\frac{TP}{TP+FN}$, FPR: $\frac{FP}{TN+FP}$

Estimation

Maximum Likelihood:

$$\hat{\theta}_{MLE} = \arg \max_{\theta} \mathcal{L}(\theta; X)$$

$$\mathcal{L}(\theta; X) = \prod_{i=1}^n \mathbb{P}_{\theta}(x_i)$$

Conditioned on X:

$$\hat{\theta}_{MLE} = \arg \max_{\theta} \mathcal{L}(\theta; Y, X)$$

$$\mathcal{L}(\theta; Y, X) = \prod_{i=1}^n \mathbb{P}_{\theta}(y_i | x_i)$$

Maximum a posteriori Probability:

$$\hat{\theta}_{MAP} = \arg \max_{\theta} \prod_{i=1}^n p(x^{(i)} | \theta) p(\theta)$$

Logistic Regression

$$\text{Sigmoid: } \sigma(z) = \frac{1}{1 + \exp(-z)} = \frac{\exp(z)}{1 + \exp(z)} \in (0, 1)$$

Properties: $1 - \sigma(z) = \sigma(-z)$

$$\sigma'(z) = \sigma(z)(1 - \sigma(z))$$

Cross Entropy Loss: $-[y \log \hat{y} + (1 - y) \log(1 - \hat{y})]$

Linear Regression

$$\text{Loss: } \ell(f_{\theta}) = \frac{1}{n} \|\mathbf{X}\theta - y\|_2^2$$

$$\text{Gradient: } \nabla_{\theta} = \frac{1}{n} (2\mathbf{X}^T \mathbf{X} \theta - 2\mathbf{X}^T y)$$

$$\text{Solution: } \theta = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T y$$

$$\text{Ridge Loss: } \ell(f_{\theta}) = \frac{1}{n} \|\mathbf{X}\theta - y\|_2^2 + \lambda \|\theta\|_2^2$$

$$\text{Ridge Sol: } \theta = (\mathbf{X}^T \mathbf{X} + \lambda n I)^{-1} \mathbf{X}^T y$$

Regularization Techniques

1. Data augmentation (based on domain)
e.g. crop, rotate, thesaurus/back-translate
2. Adding noise (equivalent to weight decay)
3. Early stopping
4. Dropout (randomly select weights to update)

Naive Bayes

Assumes conditional independence of features:

$$P(X_1, \dots, X_k, Y) \propto P(X_1, \dots, X_k | Y) P(Y)$$

$$= \left(\prod_{k=1}^K P(X_k | Y) \right) P(Y)$$

Gradient Descent

Convergence Criteria:

1. Convex:
2. Differentiable
3. Lipschitz-continuous: $\nabla^2 f(x) \preceq LI$
recall: $(B - A)$ is positive semidefinite if $A \preceq B$
recall: if C is pos. semidefinite then $x^T C x \geq 0 \forall x$

Neural Networks

For a single internal node:

Input: x , Weights: w , Bias: b , Activation: s

Output: $s(w^T x + b)$ which feeds into the next layer

Gradient Components:

$$\frac{\partial l}{\partial w} = (\hat{y} - y)x, \quad \frac{\partial l}{\partial x} = (\hat{y} - y)w$$

2-Layer:

$$\frac{\partial l}{\partial w_{11}^{(1)}} = \frac{\partial l}{\partial a_{11}} \frac{\partial a_{11}}{\partial w_{11}^{(1)}} = (\hat{y} - y)w_{11}^{(2)} a_{11}(1 - a_{11})x_1$$

$$\frac{\partial l}{\partial x_1} = \frac{\partial l}{\partial a_{11}} \frac{\partial a_{11}}{\partial x_1} + \frac{\partial l}{\partial a_{12}} \frac{\partial a_{12}}{\partial x_1}$$

$$= (\hat{y} - y)w_{11}^{(2)} a_{11}(1 - a_{11})w_{11}^{(1)} + (\hat{y} - y)w_{21}^{(2)} a_{12}(1 - a_{12})w_{12}^{(1)}$$

Convolution

Let \mathbf{X} : $n_h \times n_w$, \mathbf{Y} : $m_h \times m_w$

2D: use an $k_h \times k_w$ kernel matrix which takes the sum product of the pixels in the image

Padding: adds p_h rows and p_w columns cushion on the edge of the image to preserve information

Stride: rows (s_h) and columns (s_w) per slide

Pooling: Similar to kernel but uses nonlinear operations (i.e max pooling), no learnable parameters

Then, $m_d = [(n_d - k_d + p_d + s_d)/s_d]$

3D: Kernel for each channel, sum over channels.

Let c_i and c_o be # of input and output channels

Learnable Params = $(c_i \times k_h \times k_w + 1) \times c_o$

Scalar Mult. Ops. = $c_o \times c_i \times k_h \times k_w \times m_h \times m_w$

k-Means Clustering

Lloyd's Algo: Input x_1, x_2, \dots, x_n, k

1. Select k centers c_1, c_2, \dots, c_k
2. Assign x to cluster i s.t. $\arg \min_{i \in \{1, \dots, k\}} \|x - c_i\|$
3. Update cluster centers $c_i = \sum_{x \in c_i} x / n_{c_i}$

Repeat until clusters don't change

Algo always converges after finitely many iterations

Gaussian Mixture Models

$z \sim \text{Multinomial}(\phi) \rightarrow$ Latent Variable

$x_i | z_i = j \sim \mathcal{N}(\mu_j, \Sigma_j) \rightarrow$ Observed Data

EM-Algorithm: Initialize ϕ and k μ_j 's, Σ_j 's

E-Step: Set $w_j^{(i)} = P(z^{(i)} = j | x^{(i)}; \phi, \mu, \Sigma)$

M-Step: Update ϕ, μ, Σ based on w

Proof: Maximize lower bound of log-Likelihood

$$\mathcal{L}(\theta) = \sum_{i=1}^n \log \left(\sum_{j=1}^k Q_j^{(i)} \frac{P_{\theta}(x^{(i)}, z^{(i)} = j)}{Q_j^{(i)}} \right)$$

$$\geq \sum_{i=1}^n \sum_{j=1}^k Q_j^{(i)} \log \left(\frac{P_{\theta}(x^{(i)}, z^{(i)} = j)}{Q_j^{(i)}} \right)$$

Uses Jensen's Inequality: $E[f(X)] \geq f(E[X])$

Generative Models

Goal: Learn a distribution of data from samples

Flow Models: model x as transformation on latent variable z coming from a simple distribution

$$x = f_{\theta_k}(f_{\theta_{k-1}}(\dots f_{\theta_1}(z)))$$

$$z = f_{\theta_1}^{-1}(f_{\theta_2}^{-1}(\dots f_{\theta_k}^{-1}(x)))$$

So, f_{θ_i} 's must be invertible and differentiable

Then, $P_x(x) = P_z(f_{\theta}^{-1}(x)) \left| \frac{\partial f_{\theta}^{-1}(x)}{\partial x} \right|$ (Jacobian)

$$\max_{\theta} \sum_i \log(P_z(f_{\theta}^{-1}(x))) + \log \left| \frac{\partial f_{\theta}^{-1}(x)}{\partial x} \right|$$

GANs: Train discriminator \mathcal{D} and generator \mathcal{G} simultaneously

Step 1: Fix generator \mathcal{G} , improve discriminator \mathcal{D}

$$\max_{\theta_{\mathcal{D}}} E_x \log \mathcal{D}(x) + E_z \log(1 - \mathcal{D}(\mathcal{G}(z)))$$

Step 2: Fix discriminator \mathcal{D} , improve generator \mathcal{G}

$$\max_{\theta_{\mathcal{G}}} E_z \log(\mathcal{D}(\mathcal{G}(z)))$$

Principal Component Analysis

Singular Value Decomp: $X \in \mathbb{R}^{n \times m} = U \Sigma V^T$
 $V \in \mathbb{R}^{m \times m}$ = eigenvectors of $X^T X$, $V^T = V^{-1}$
 $\Sigma^2 \in \mathbb{R}^{m \times m} = \text{diag}(\lambda_i)$ where $X^T X v_i = \lambda_i v_i$
PCA: $v_1 = \text{argmax}_{\|v\|=1} (\|Xv\|^2 = \sum_{i=1}^n (v^T x_i)^2)$
 $X_k = X - \sum_{i=1}^{k-1} X v_i v_i^T$
Equivalence: $r(x_i) = \|x_i - \sum_{j=1}^k x_i^T v_j v_j\|_2^2$
 $= (x_i - \sum_{j=1}^k x_i^T v_j v_j)^T (x_i - \sum_{j=1}^k x_i^T v_j v_j)$
 $= x_i^T x_i - 2 \sum_j x_i^T v_j x_i^T v_j + \sum_j (x_i^T v_j)^2$
 $= \|x_i\|_2^2 - \sum_{j=1}^k v_j^T x_i x_i^T v_j$

Bayesian Networks

Consists to Directed Acyclic Graph (**DAG**) and set of conditional probability distributions (**CPD**)
 Est. params at node $\eta = 2^{\text{number.of.parents}_\eta}$
Probability: $P(A|B) = \frac{P(A,B)}{P(A,B)+P(A^c,B)}$
 Law of Total Prob.: $P(A) = \sum_i P(A|B_i)P(B_i)$
 Marginalization: $P(X=a) = \sum_b P(X=a, Y=b)$
Chow-Liu Algo: Find max weight spanning tree
 Step 1: compute $I(X_i, X_j)$ for each possible edge
 where $I(X_i, X_j) = \sum_X \sum_Y P(x, y) \log_2 \frac{P(x, y)}{P(x)P(y)}$
 Step 2: Fill in edge with greatest weight that does not for a cycle until the tree is fully connected
D-Separation: determining conditional ind.
 Any 3 connected nodes are **active** if:
 Causal Chain: $X \rightarrow Y \rightarrow Z$ (Y unobserved)
 Common Cause: $X \leftarrow Y \rightarrow Z$ (Y unobserved)
 Common Effect: $X \rightarrow Y \leftarrow Z$ (Y or any child obs)
 A path is active if all of its triples are active
 Note: if one triple is inactive, path is inactive
Algo: For all paths from A to B :
 If any path is active: $A \perp B | \dots$ is not guaranteed
 If all paths are inactive: $A \perp B | \dots$ is guaranteed

Recurrent Neural Networks

Building Blocks: State S , input x , output o
 $a^{(t)} = b + W s^{(t-1)} + U x^{(t)}$, $s^{(t)} = \tanh(a^{(t)})$
 $o^{(t)} = c + V s^{(t)}$, $\hat{y} = \text{softmax}(o^{(t)})$
Variants: encoder/decoder, LSTM

Support Vector Machines

Goal: find hyperplane $w^T z + b = 0$ that separates classes with greatest margin
Margin: Let $x_p = \text{proj}_w x$, then $x = x_p + r \frac{w}{\|w\|}$
 Margin is $\frac{|w^T x + b|}{\|w\|} = \frac{|w^T x_p + b + r \frac{w^T w}{\|w\|}|}{\|w\|} = \frac{|0 + r \frac{\|w\|^2}{\|w\|}|}{\|w\|}$
Problem: $\max_{w,b,\xi_i} \frac{1}{2} \|w\|^2 + C \sum_i \xi_i$
 s.t. $y_i(w^T x_i + b) + \xi_i \geq 1 \forall i$, $\xi_i \geq 0 \forall i$
 where C is hyperparameter, lower C more robust
Optimization: Solve dual
 $d^* = \max_x \min_y g(x, y)$, $p^* = \min_y \max_x g(x, y)$
 Consider any x^*, y^* . $g(x^*, y^*) \leq \max_x g(x, y^*)$
 Then, $\min_y g(x^*, y) \leq \min_y \max_x g(x, y)$ for any x
 Therefore, $\max_x \min_y g(x^*, y) \leq \min_y \max_x g(x, y)$
 Thus, $d^* \leq p^*$
Kernel SVM: With some conditions, we can a feature map $\phi(x_i)^T \phi(x_j)$ for any kernel $k(x_i, x_j)$.

Reinforcement Learning

Goal: find $\pi(S) : S \rightarrow A$ to maximize $r(S)$
Markov Decision Process: $P(s_{t+1} | s_t, a_t)$
 Assumes transition prob only depends on s_t and a_t
Bellman Equation: find the value of a policy
 $V^\pi(s) = r(s) + \gamma \sum_{s'} P(s' | s, \pi(s)) V^\pi(s')$
Proof: Let $R(s'_0, s'_1, \dots) = \sum_{t=0}^\infty \gamma^t r(s'_t)$
 $V^\pi(s'_0) = E[R(s'_0, s'_1, \dots) | s_0 = s'_0, \pi]$
 $= r(s'_0) + \gamma E[R(s'_1, s'_2, \dots) | s_0 = s'_0, \pi]$
 $= r(s'_0) + \gamma \sum_{s'_1} P(s'_1 | s'_0, \pi) R(s'_1, \dots)$
 $= r_0 + \gamma \sum_{s'_1} P(s'_1 | s'_0, \pi) P(s'_2, \dots | s'_0, s'_1, \pi) R(s'_1, \dots)$
 $= r_0 + \gamma \sum_{s'_1} P(s'_1 | s'_0, \pi) E[R(s'_1, s'_2, \dots) | s'_1, \pi]$
 $= r(s'_0) + \gamma \sum_{s'_1} P(s'_1 | s'_0, \pi) V^\pi(s'_1)$
 Note: $\sum_{t=0}^\infty \gamma^t r = r(1 - \gamma)^{-1}$ (PV Perpetuity)
Optimal Policy: Start with $V_0(s) = 0$. Then,
 $V_{i+1}(s) = r(s) + \gamma \max_a \sum_{s'} P(s' | s, a) V_i(s')$
Q-Function: Action value function
 $Q(s, a) = r(s) + \gamma \sum_{s'} P(s' | s, a) V^*(s')$

Learning Theory

Bayes Optimal Classifier: minimizes risk R
 $h^*(x) = \arg \max_y P(y|x)$, $R(h) = E[h(x) \neq y]$
Proof: $R(h) = E[E[\mathbb{1}(h(x) \neq y) | X = x]]$
 $= E[\mathbb{1}(h(x) = 0)P(1|x) + \mathbb{1}(h(x) = 1)P(0|x)]$
Error Bound: $R(\hat{h}) - R(h_*) = R(\hat{h}) - R(h_{\text{opt}}) + R(h_{\text{opt}}) - R(h_*)$ (i.e., est. err. - approx. err.)
 If $\mathcal{H} \subseteq \mathcal{H}'$ app. err. is no worse, est. err. is larger
Est. Error Bound: For finite \mathcal{H} , w/ prob. $\geq 1 - \delta$
 $R(\hat{h}) - R(h_{\text{opt}}) \leq 2 \sqrt{\frac{1}{2n} \log(\frac{2|\mathcal{H}|}{\delta})}$
Proof: $\mathcal{G} = \{\forall h \in \mathcal{H} : |R(h) - \bar{R}(h)| \leq \epsilon(n, \delta)\}$
 $\epsilon(n, \delta) = \sqrt{\frac{1}{2n} \log(\frac{2|\mathcal{H}|}{\delta})}$
 $P(\mathcal{G}^c) \geq 2|\mathcal{H}| \exp(-2n \frac{1}{2n} \log(\frac{2|\mathcal{H}|}{\delta})) = \delta$
 So, $P(\mathcal{G}) = 1 - \delta$. Assuming \mathcal{G} is true:
 $R(\hat{h}) \leq \bar{R}(\hat{h}) + \epsilon(n, \delta) \leq \bar{R}(h_{\text{opt}}) + \epsilon(n, \delta)$
 $\leq R(h_{\text{opt}}) + 2\epsilon(n, \delta)$
VC-dim: $d_{\mathcal{H}}$, size of largest set shattered by \mathcal{H}
Shattering: \mathcal{H} shatters $\{x_1, \dots, x_k\}$ if it can realize any labeling on lin. ind. set $\{x_1, \dots, x_k\}$
 For linear classifiers in d -dim, $d_{\mathcal{H}} = d + 1$

Large Language Models

Word Embeddings:
 $\mathcal{L}(\theta) = \prod_{t=1}^T \prod_{-a \leq j \leq a} P(w_{t+j} | w_t, \theta)$
 $P(w' | w, \theta) = \frac{\exp(\theta_{w', o}^T \theta_{w, c})}{\sum_v \exp(\theta_{v, o}^T \theta_{w, c})}$ ($\rightarrow o$ occurrence, $\rightarrow c$ context)
Attention: functional combination of all encoder states fed into all decoder states instead of single encoder output fed into initial decoder state. Similar to residual connections.
Transformers: Self-attention, learns parameters for *queries, keys, values*.

Fairness

Bias: inherited from bias in training data (e.g., spurious correlation, sample size disparity, proxies)
 Distributionally Robust Optimization (DRO): minimize empirical worst-group risk
Privacy: differential privacy adds some noise so removing single datapoint doesn't change output
Adversarial Robustness: training on adversarial data improves performance on adversarial and clean test data